


第11讲 图像生成模型 (上)

神经网络与深度学习 2026

信息科学技术学院 吴瀚霖

hlwu@bfsu.edu.cn




2021

2022

2024

2015



DALL·E (OpenAI)




DALL·E 2 (OpenAI)



Imagen (Google)



Latent Diffusion Models (Stability AI)



可灵 klingai.com



Sora (OpenAI)

2025

2026

1 更清晰的文字生成 (Readable Text Rendering)



海报、招牌、界面、杂志、包装等清晰可读

2 多模态统一生成 (Multimodal Generation)



图像、视频、音频、3D 统一理解与生成
GPT-4o / Gemini 2.0 等原生多模态模型

3 可控图像编辑 (Controllable Image Editing)



指令式编辑 (Instruction-based Editing)



精准控制内容、风格、颜色和主体一致性

4 长视频生成 (Long Video Generation)



长时稳定生成 (分钟级以上连续视频)



高质量、连贯的长视频叙事

5 世界模型与视频理解生成 (World Models for Video)



交互式场景生成与预测

参考这四张PPT，逐张
配上voiceover，正常语
速，男声。



AI





生成模型的特点

- **一对多**（甚至无穷）的映射
- 预测结果的**维数**通常远高于输入
- 一些预测的**可能性**比其他预测更高
- 训练数据**不可能穷尽**所有的解



DALL-E 2 (OpenAI)



Imagen (Google)

GAN

2014.06:

- Goodfellow, “*Generative Adversarial Nets*”.

What are some recent and potentially upcoming breakthroughs in unsupervised learning?



Yann LeCun, Director of AI Research at Facebook and Professor at NYU
Written Jul 29 · Upvoted by Joaquin Quiñero Candela, [Director Applied Machine Learning at Facebook](#) and Huang Xiao



Adversarial training is the coolest thing since sliced bread.

I've listed a bunch of relevant papers in a previous answer.

Expect more impressive results with this technique in the coming years.

What's missing at the moment is a good understanding of it so we can make it work reliably. It's very finicky. Sort of like ConvNet were in the 1990s, when I had the reputation of being the only person who could make them work (which wasn't true).



Goodfellow

The GAN zoo

<https://github.com/hindupuravinash/the-gan-zoo>

GAN

ACGAN

BGAN

CGAN


DCGAN

EBGAN

fGAN

GoGAN


⋮



Mihaela Rosca, Balaji Lakshminarayanan, David Warde-Farley, Shakir Mohamed, "Variational Approaches for Auto-Encoding Generative Adversarial Networks", arXiv, 2017


²We use the Greek α prefix for α -GAN, as AEGAN and most other Latin prefixes seem to have been taken
<https://deephunt.in/the-gan-zoo-79597dc8c347>.

GAN的基本思想





GAN的基本思想

通过调整 z 的不同维度，可以生成不同属性的图像



GAN的基本思想




→ 判别器 → 1.0



→ 判别器 → 0.9



→ 判别器 → 0.1




→ 判别器 → 0.1

GAN的基本思想

生成器



棕色的




枯叶蝶

蝴蝶不是棕色的




蝴蝶没有脉络




.....

判别器

GAN的基本思想



GAN的基本思想



GAN的基本思想

初始化生成器与判别器




在每次迭代中

- **STEP1**: 固定生成器 G , 更新判别器 D 的参数

损失函数:

$$-\mathbf{E}_{(x,y)}[y \log D(x) + (1 - y) \log(1 - D(x))]$$



GAN的基本思想

初始化生成器与判别器




在每次迭代中

- **STEP2**: 固定判别器 D , 更新生成器 G 的参数




生成模型在学什么

图像可以看作一个向量




生成模型在学什么



生成模型在学什么

并不是所有的点都对应自然图像

好的图像概率高，差的图像概率低



核心思想： 生成模型不是记住某一张图，而是在学习 “什么样的图像更可能出现”。




训练 = 拟合分布；



生成 = 从分布中采样。


生成模型的本质是建模一个概率分布




* 这两幅图像是使用Grok生成的

如何建模概率分布？

一个映射就搞定



GAN的优化目标




$$G^* = \arg \min \text{KL}(\mathbf{P}_G || \mathbf{P}_{\text{data}})$$

某种衡量分布距离的函数

GAN的优化目标

判别器 D




判别器的作用就是帮我们衡量分布间的距离




北京外国语大学
BEIJING FOREIGN STUDIES UNIVERSITY

VAE

自编码器




自编码器




反卷积

低维特征映射到高维特征






自编码器可以降维



结论：降维不是目的，识别并保留数据的主要结构，才是核心。

什么叫保留数据的结构




什么叫有意义的表示





自编码器作为生成模型


解码器看起来能生成图像




自编码器作为生成模型




原始的自编码器




加点噪声





加点噪声（等价的形式）



变分自编码器的实现



正态分布有要求吗




每一个正态分布都要向标准正态分布看齐

变分自编码器

损失函数

- 第一部分：重建损失
- 第二部分：正则化项



$$\sum_i (\exp(\tilde{\sigma}_i) - (1 + \tilde{\sigma}_i) + \mu_i^2)$$

$\tilde{\sigma}_i = \log \sigma^2$ 应在0附近;
隐编码的分布不要退化为单点分布

隐编码不要太分散 (应在0附近)

正则化项为 $\text{KL}(N(\mu_i, \sigma_i^2) || N(0,1))$.