



第12讲 图像生成模型（下）

神经网络与深度学习 2026

信息科学技术学院 吴瀚霖

hlwu@bfsu.edu.cn

生成模型在学什么



核心思想： 生成模型不是记住某一张图，而是在学习 “什么样的图像更可能出现”。




训练 = 拟合分布；




生成 = 从分布中采样。

如何建模概率分布？


一个映射就搞定



GAN的基本思想



GAN的优化目标




$$G^* = \arg \min \text{KL}(\mathbf{P}_G || \mathbf{P}_{\text{data}})$$


某种衡量分布距离的函数

GAN的优化目标

判别器 D



判别器的作用就是帮我们衡量分布间的距离




2021

2022

2024

2015



DALL-E (OpenAI)




DALL-E 2 (OpenAI)



Imagen (Google)



Latent Diffusion Models (Stability AI)



可灵 klingai.com





Sora (OpenAI)

扩散模型的基本原理

从噪声到图像的映射

逆向过程





The figure you see in the finished sculpture is already there inside the marble.
My work is simply to remove the excess material.



1. The block as it is
The sculpture exists within the marble.



2. Removing the excess
Chiseling away the material that is not part of the figure.





3. Refining the form
Continuing to reveal the details that are already there.




4. The finished sculpture
The figure is revealed—it was always inside.

“ I do not create the sculpture; I set it free. — Michelangelo

每一步如何去噪



每一步如何去噪

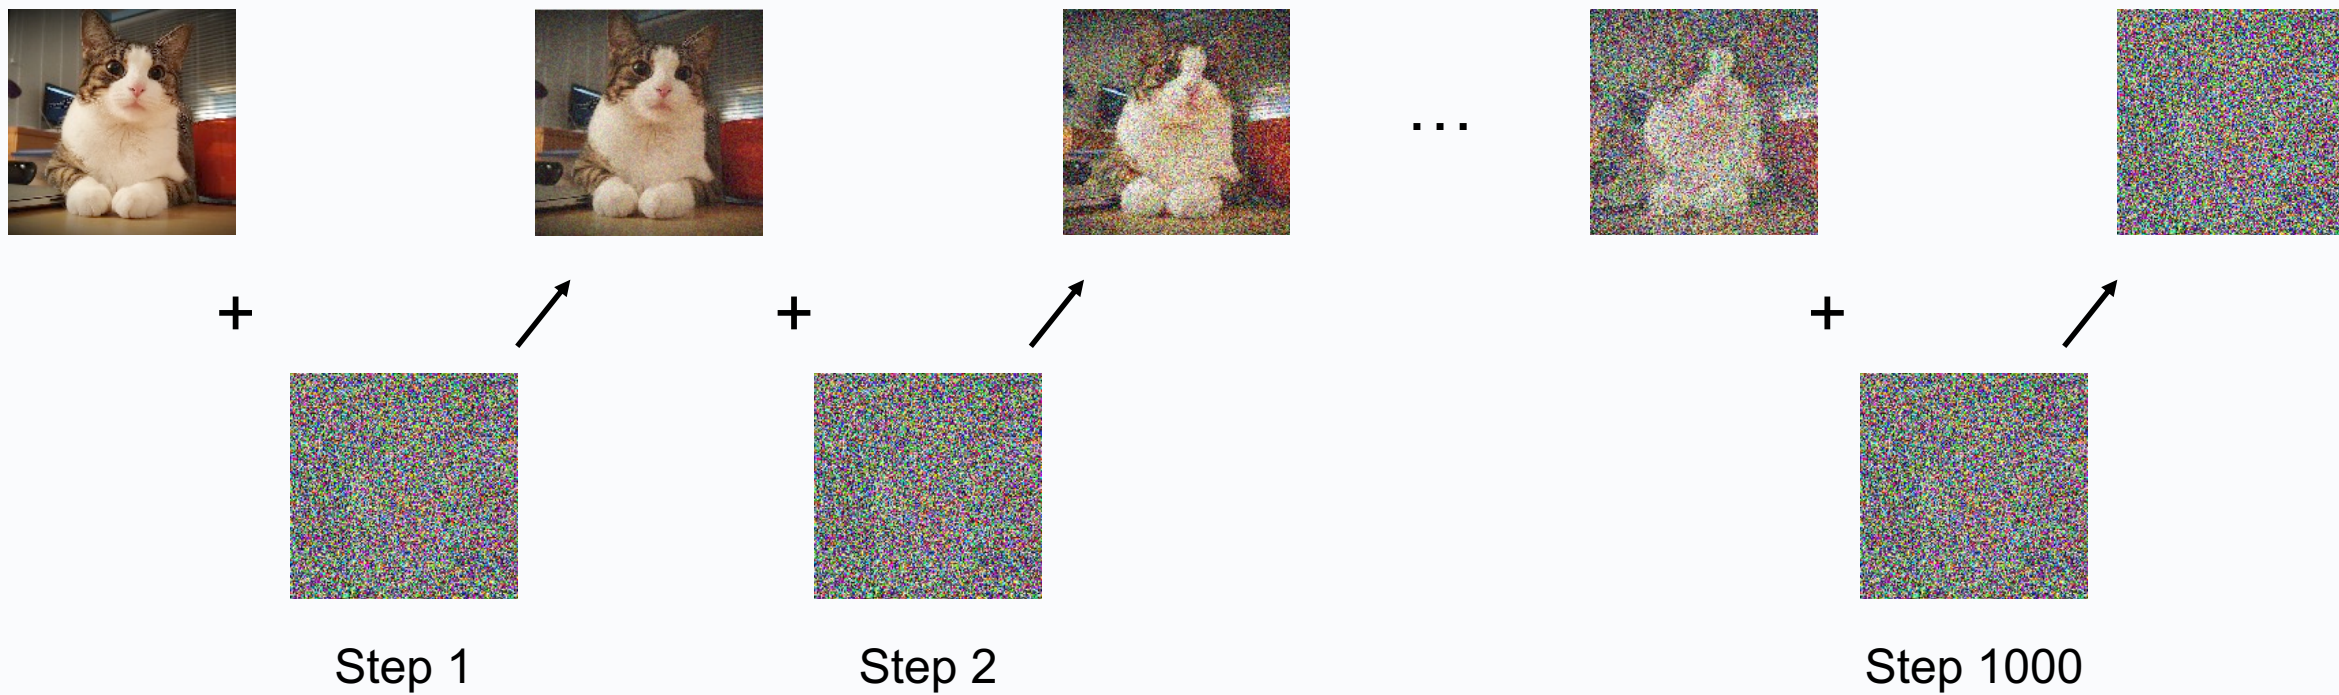


如何训练噪声预测器

Creating noise from data is easy.


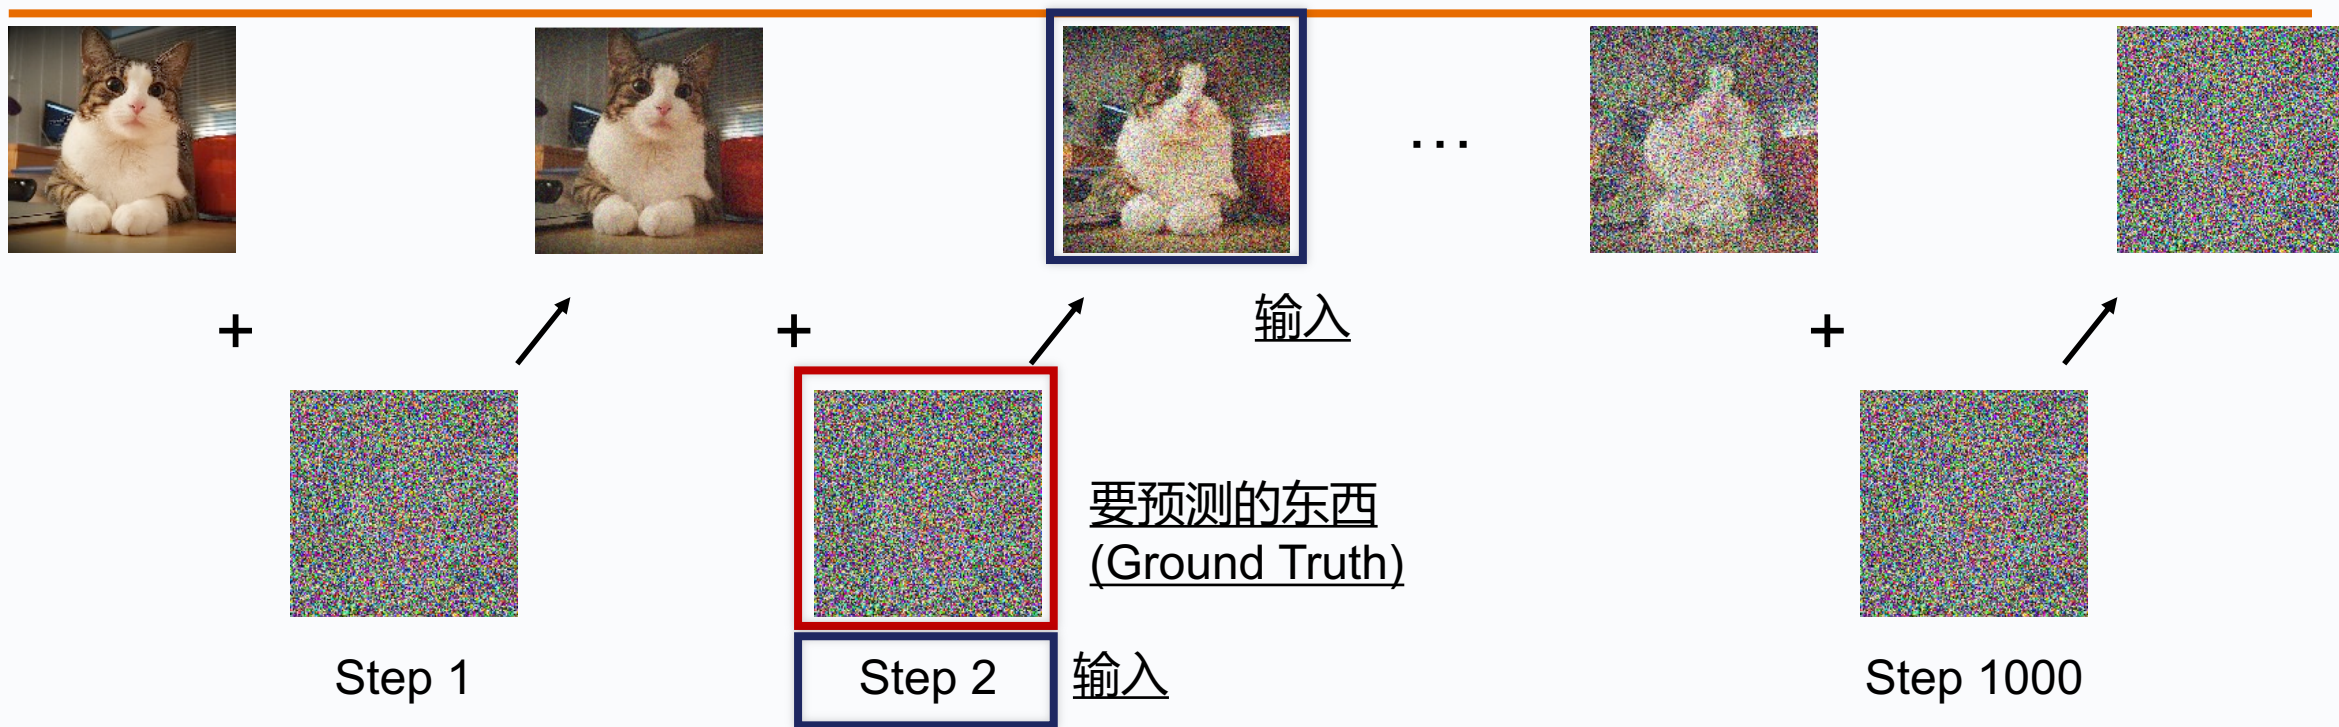
Creating data from noise is generative modeling.

正向过程

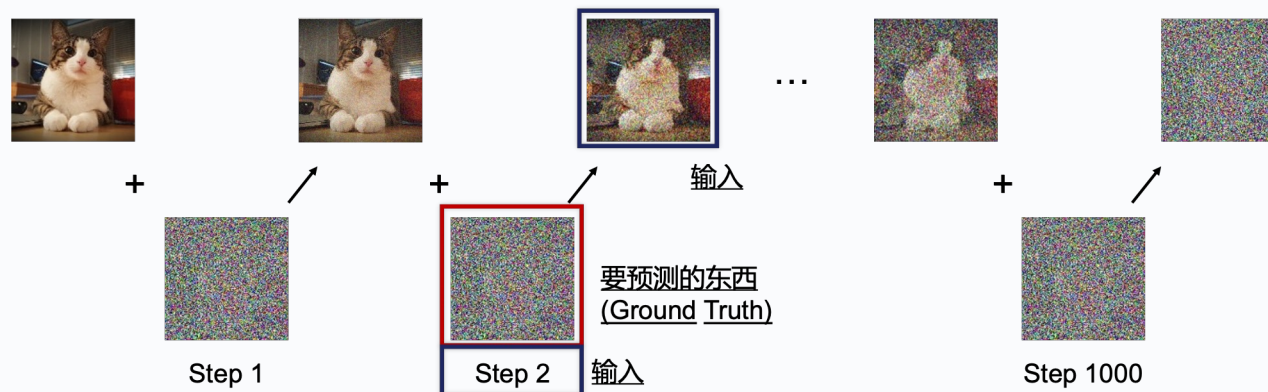


扩散过程（正向过程）

训练目标



DDPM



Denoising Diffusion Probabilistic Models, 2020

Algorithm 1 Training


- 1: **repeat**
- 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
- 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
- 4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 5: Take gradient descent step on

$$\nabla_{\theta} \left\| \epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2$$
- 6: **until** converged

Algorithm 2 Sampling

- 1: $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 2: **for** $t = T, \dots, 1$ **do**
- 3: $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ if $t > 1$, else $\mathbf{z} = \mathbf{0}$
- 4: $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$
- 5: **end for**
- 6: **return** $\mathbf{x}_0$

为什么训练目标长这样？




想象中:

$$X_t = X_{t-1} + \epsilon$$

实际上:

$$X_t = \sqrt{1 - \beta_t} X_{t-1} + \sqrt{\beta_t} \epsilon$$

为什么训练目标长这样？



可以跳着走

$$X_t = \sqrt{1 - \beta_t} X_{t-1} + \sqrt{\beta_t} \epsilon$$


$$X_{t-1} = \sqrt{1 - \beta_{t-1}} X_{t-2} + \sqrt{\beta_{t-1}} \epsilon$$

$$X_t = \sqrt{(1 - \beta_t)(1 - \beta_{t-1}) \cdots (1 - \beta_1)} X_0 + \sqrt{1 - (1 - \beta_t)(1 - \beta_{t-1}) \cdots (1 - \beta_1)} \epsilon$$

$$X_t = \sqrt{\alpha_t} X_{t-1} + \sqrt{1 - \alpha_t} \epsilon$$

为什么训练目标长这样？

实际的训练



Algorithm 1 Training

- 1: **repeat**
- 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
- 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
- 4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 5: Take gradient descent step on
 $\nabla_{\theta} \|\epsilon - \epsilon_{\theta}(\sqrt{\alpha_t} \mathbf{x}_0 + \sqrt{1 - \alpha_t} \epsilon, t)\|^2$
- 6: **until** converged

模型：噪声预测器

为什么推理长这样？

$$\begin{aligned}
 \log p(\mathbf{x}) &\geq \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T) \prod_{t=1}^T p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{\prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1})} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T) p_{\theta}(\mathbf{x}_0|\mathbf{x}_1) \prod_{t=2}^T p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_1|\mathbf{x}_0) \prod_{t=2}^T q(\mathbf{x}_t|\mathbf{x}_{t-1})} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T) p_{\theta}(\mathbf{x}_0|\mathbf{x}_1) \prod_{t=2}^T p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_1|\mathbf{x}_0) \prod_{t=2}^T q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0)} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p_{\theta}(\mathbf{x}_T) p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \log \prod_{t=2}^T \frac{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0)} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T) p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \log \prod_{t=2}^T \frac{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{\frac{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) q(\mathbf{x}_t|\mathbf{x}_0)}{q(\mathbf{x}_{t-1}|\mathbf{x}_0)}} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T) p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \log \prod_{t=2}^T \frac{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_t|\mathbf{x}_0)} + \log \prod_{t=2}^T \frac{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T) p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)}{q(\mathbf{x}_1|\mathbf{x}_0)} + \sum_{t=2}^T \log \frac{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} [\log p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)] + \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)}{q(\mathbf{x}_T|\mathbf{x}_0)} \right] + \sum_{t=2}^T \mathbb{E}_{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[\log \frac{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
 &= \mathbb{E}_{q(\mathbf{x}_1|\mathbf{x}_0)} [\log p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)] + \mathbb{E}_{q(\mathbf{x}_T|\mathbf{x}_0)} \left[\log \frac{p(\mathbf{x}_T)}{q(\mathbf{x}_T|\mathbf{x}_0)} \right] + \sum_{t=2}^T \mathbb{E}_{q(\mathbf{x}_t, \mathbf{x}_{t-1}|\mathbf{x}_0)} \left[\log \frac{p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)}{q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)} \right] \\
 &= \underbrace{\mathbb{E}_{q(\mathbf{x}_1|\mathbf{x}_0)} [\log p_{\theta}(\mathbf{x}_0|\mathbf{x}_1)]}_{\text{reconstruction term}} - \underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \parallel p(\mathbf{x}_T))}_{\text{prior matching term}} - \sum_{t=2}^T \underbrace{\mathbb{E}_{q(\mathbf{x}_t|\mathbf{x}_0)} [D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \parallel p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t))]}_{\text{denoising matching term}}
 \end{aligned}$$

$$\begin{aligned}
 q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) &= \frac{q(\mathbf{x}_t|\mathbf{x}_{t-1}, \mathbf{x}_0) q(\mathbf{x}_{t-1}|\mathbf{x}_0)}{q(\mathbf{x}_t|\mathbf{x}_0)} \\
 &= \frac{\mathcal{N}(\mathbf{x}_t; \sqrt{\alpha_t} \mathbf{x}_{t-1}, (1-\alpha_t) \mathbf{I}) \mathcal{N}(\mathbf{x}_{t-1}; \sqrt{\alpha_{t-1}} \mathbf{x}_0, (1-\alpha_{t-1}) \mathbf{I})}{\mathcal{N}(\mathbf{x}_t; \sqrt{\alpha_t} \mathbf{x}_0, (1-\alpha_t) \mathbf{I})} \\
 &\propto \exp \left\{ - \left[\frac{(\mathbf{x}_t - \sqrt{\alpha_t} \mathbf{x}_{t-1})^2}{2(1-\alpha_t)} + \frac{(\mathbf{x}_{t-1} - \sqrt{\alpha_{t-1}} \mathbf{x}_0)^2}{2(1-\alpha_{t-1})} - \frac{(\mathbf{x}_t - \sqrt{\alpha_t} \mathbf{x}_0)^2}{2(1-\alpha_t)} \right] \right\} \\
 &= \exp \left\{ - \frac{1}{2} \left[\frac{(\mathbf{x}_t - \sqrt{\alpha_t} \mathbf{x}_{t-1})^2}{1-\alpha_t} + \frac{(\mathbf{x}_{t-1} - \sqrt{\alpha_{t-1}} \mathbf{x}_0)^2}{1-\alpha_{t-1}} - \frac{(\mathbf{x}_t - \sqrt{\alpha_t} \mathbf{x}_0)^2}{1-\alpha_t} \right] \right\} \\
 &= \exp \left\{ - \frac{1}{2} \left[\frac{(-2\sqrt{\alpha_t} \mathbf{x}_t \mathbf{x}_{t-1} + \alpha_t \mathbf{x}_{t-1}^2)}{1-\alpha_t} + \frac{(\mathbf{x}_{t-1}^2 - 2\sqrt{\alpha_{t-1}} \mathbf{x}_{t-1} \mathbf{x}_0)}{1-\alpha_{t-1}} + C(\mathbf{x}_t, \mathbf{x}_0) \right] \right\} \\
 &\propto \exp \left\{ - \frac{1}{2} \left[- \frac{2\sqrt{\alpha_t} \mathbf{x}_t \mathbf{x}_{t-1}}{1-\alpha_t} + \frac{\alpha_t \mathbf{x}_{t-1}^2}{1-\alpha_t} + \frac{\mathbf{x}_{t-1}^2}{1-\alpha_{t-1}} - \frac{2\sqrt{\alpha_{t-1}} \mathbf{x}_{t-1} \mathbf{x}_0}{1-\alpha_{t-1}} \right] \right\} \\
 &= \exp \left\{ - \frac{1}{2} \left[\left(\frac{\alpha_t}{1-\alpha_t} + \frac{1}{1-\alpha_{t-1}} \right) \mathbf{x}_{t-1}^2 - 2 \left(\frac{\sqrt{\alpha_t} \mathbf{x}_t}{1-\alpha_t} + \frac{\sqrt{\alpha_{t-1}} \mathbf{x}_0}{1-\alpha_{t-1}} \right) \mathbf{x}_{t-1} \right] \right\} \\
 &= \exp \left\{ - \frac{1}{2} \left[\frac{\alpha_t(1-\alpha_{t-1}) + 1 - \alpha_t}{(1-\alpha_t)(1-\alpha_{t-1})} \mathbf{x}_{t-1}^2 - 2 \left(\frac{\sqrt{\alpha_t} \mathbf{x}_t}{1-\alpha_t} + \frac{\sqrt{\alpha_{t-1}} \mathbf{x}_0}{1-\alpha_{t-1}} \right) \mathbf{x}_{t-1} \right] \right\} \\
 &= \exp \left\{ - \frac{1}{2} \left[\frac{\alpha_t - \alpha_t + 1 - \alpha_t}{(1-\alpha_t)(1-\alpha_{t-1})} \mathbf{x}_{t-1}^2 - 2 \left(\frac{\sqrt{\alpha_t} \mathbf{x}_t}{1-\alpha_t} + \frac{\sqrt{\alpha_{t-1}} \mathbf{x}_0}{1-\alpha_{t-1}} \right) \mathbf{x}_{t-1} \right] \right\} \\
 &= \exp \left\{ - \frac{1}{2} \left[\frac{1 - \alpha_t}{(1-\alpha_t)(1-\alpha_{t-1})} \mathbf{x}_{t-1}^2 - 2 \left(\frac{\sqrt{\alpha_t} \mathbf{x}_t}{1-\alpha_t} + \frac{\sqrt{\alpha_{t-1}} \mathbf{x}_0}{1-\alpha_{t-1}} \right) \mathbf{x}_{t-1} \right] \right\} \\
 &= \exp \left\{ - \frac{1}{2} \left(\frac{1 - \alpha_t}{(1-\alpha_t)(1-\alpha_{t-1})} \right) \left[\mathbf{x}_{t-1}^2 - 2 \frac{\left(\frac{\sqrt{\alpha_t} \mathbf{x}_t}{1-\alpha_t} + \frac{\sqrt{\alpha_{t-1}} \mathbf{x}_0}{1-\alpha_{t-1}} \right)}{\frac{1-\alpha_t}{(1-\alpha_t)(1-\alpha_{t-1})}} \mathbf{x}_{t-1} \right] \right\} \\
 &= \exp \left\{ - \frac{1}{2} \left(\frac{1 - \alpha_t}{(1-\alpha_t)(1-\alpha_{t-1})} \right) \left[\mathbf{x}_{t-1}^2 - 2 \frac{\left(\frac{\sqrt{\alpha_t} \mathbf{x}_t}{1-\alpha_t} + \frac{\sqrt{\alpha_{t-1}} \mathbf{x}_0}{1-\alpha_{t-1}} \right) (1-\alpha_t)(1-\alpha_{t-1})}{1-\alpha_t} \mathbf{x}_{t-1} \right] \right\} \\
 &= \exp \left\{ - \frac{1}{2} \left(\frac{1}{(1-\alpha_t)(1-\alpha_{t-1})} \right) \left[\mathbf{x}_{t-1}^2 - 2 \frac{\sqrt{\alpha_t}(1-\alpha_{t-1}) \mathbf{x}_t + \sqrt{\alpha_{t-1}}(1-\alpha_t) \mathbf{x}_0}{1-\alpha_t} \mathbf{x}_{t-1} \right] \right\} \\
 &\propto \mathcal{N}(\mathbf{x}_{t-1}; \underbrace{\frac{\sqrt{\alpha_t}(1-\alpha_{t-1}) \mathbf{x}_t + \sqrt{\alpha_{t-1}}(1-\alpha_t) \mathbf{x}_0}{1-\alpha_t}}_{\mu_q(\mathbf{x}_t, \mathbf{x}_0)}, \underbrace{\frac{(1-\alpha_t)(1-\alpha_{t-1})}{1-\alpha_t}}_{\Sigma_q(t)} \mathbf{I})
 \end{aligned}$$

Algorithm 2 Sampling

- 1: $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 2: **for** $t = T, \dots, 1$ **do**
- 3: $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ if $t > 1$, else $\mathbf{z} = \mathbf{0}$
- 4: $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\alpha_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$
- 5: **end for**
- 6: **return** $\mathbf{x}_0$

DDPM生成的图片示例



Figure 3: LSUN Church samples. FID=7.89





Figure 4: LSUN Bedroom samples. FID=4.90

FID: 评价图像生成的常用指标




red points: real images

blue points: generated images

FID = Fréchet distance
between the two **Gaussians** ???

Smaller is better

A lot of samples
is needed.






北京外国语大学
BEIJING FOREIGN STUDIES UNIVERSITY


Demo

图像生成的通用套路




中间产物是什么？

数据内在结构位于一个低维流形上。



潜在扩散模型 (LDM)



文生图模型的训练数据集




训练数据集



一只在奔跑的狗



雪地里的猫




阳光下的猫



沙滩上的狗

训练潜在扩散模型 (LDM)




LDM的文生图

1 一只猫趴在桌子上



2



Stable Diffusion的结果

Text-to-Image Synthesis on LAION. 1.45B Model.

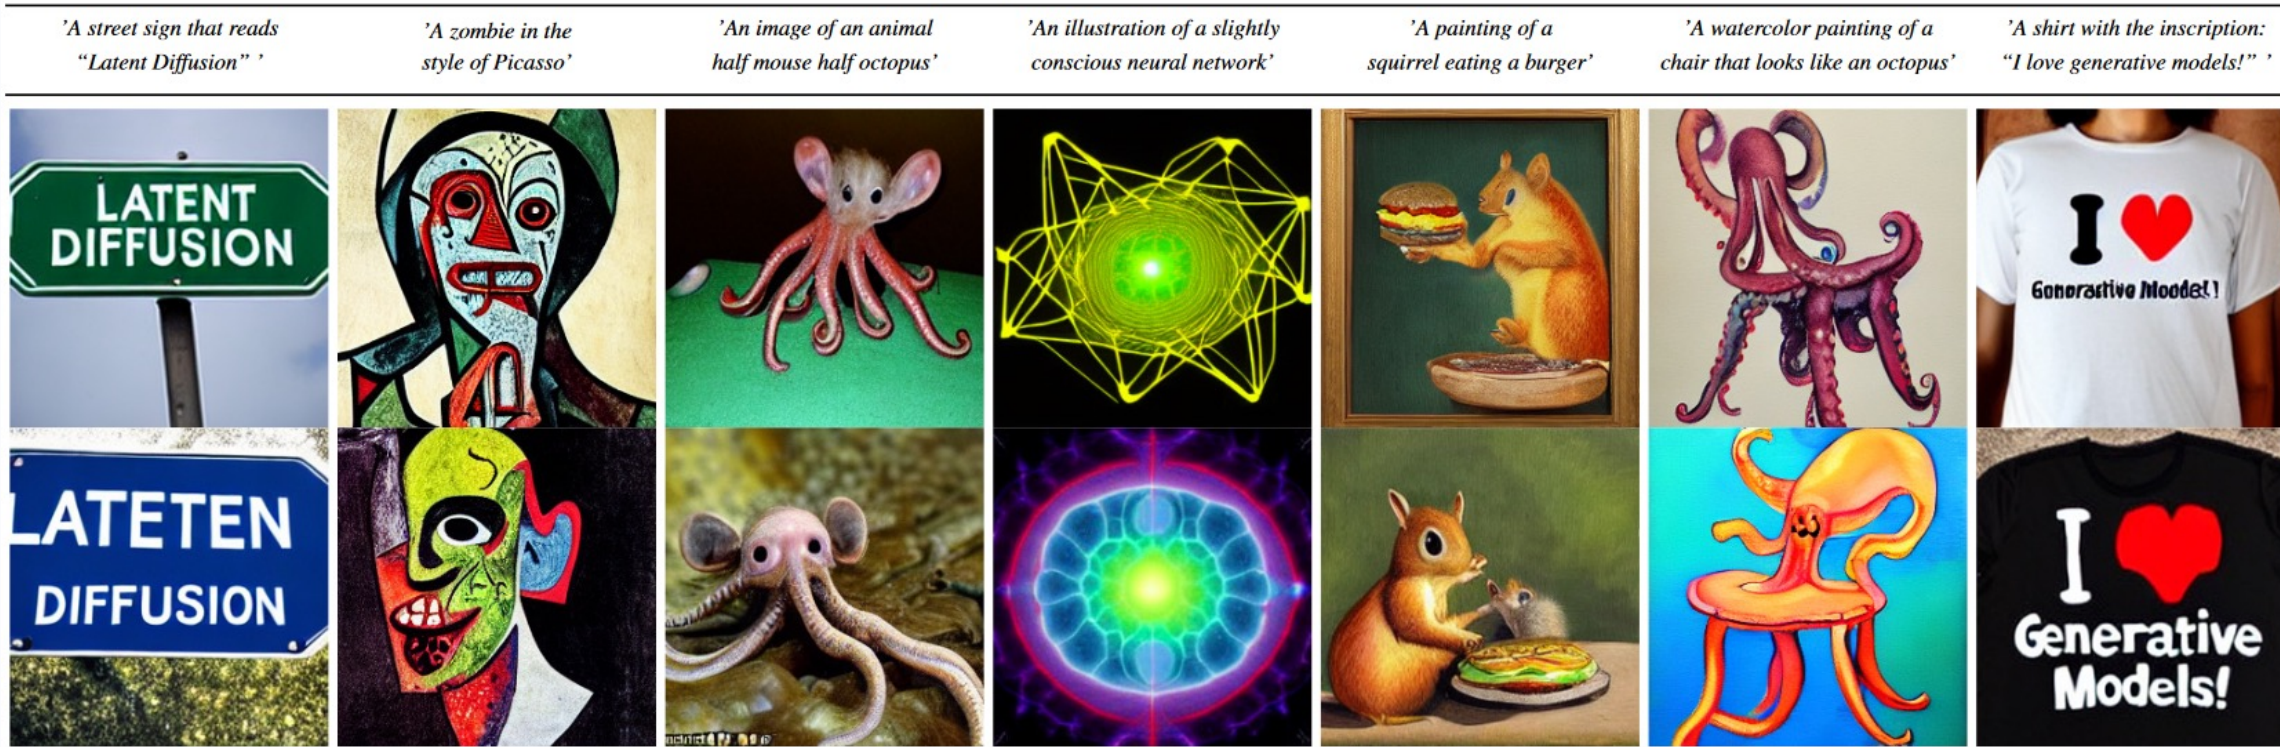


Figure 5. Samples for user-defined text prompts from our model for text-to-image synthesis, *LDM-8 (KL)*, which was trained on the LAION [78] database. Samples generated with 200 DDIM steps and $\eta = 1.0$. We use unconditional guidance [32] with $s = 10.0$.

Laion

<https://laion.ai/blog/laion-5b/>

5.85B

Backend url:

Index:

Search: 🔍 📷 ⬇

[Clip retrieval](#) works by converting the text query to a CLIP embedding, then using that embedding to query a knn index of clip image embedddings

Display captions ☒
Display full captions ☐
Display similarities ☐
Safe mode ☒
Hide duplicate urls ☒
Hide (near) duplicate images ☒
Search over
Search with multilingual clip ☐


french cat


french cat


How to tell if your feline is french. He wears a b...


イケメン猫モデル「トキ・ナンタケット」がカッコいい - NAVER まとめ


Hilarious pics of funny cats! funnycatsgif.com


Hipster cat


網友挑戰「加幾筆畫出最創意貓咪圖片」，笑到岔氣之後我也手


cat in a suit Georgian sells tomatoes

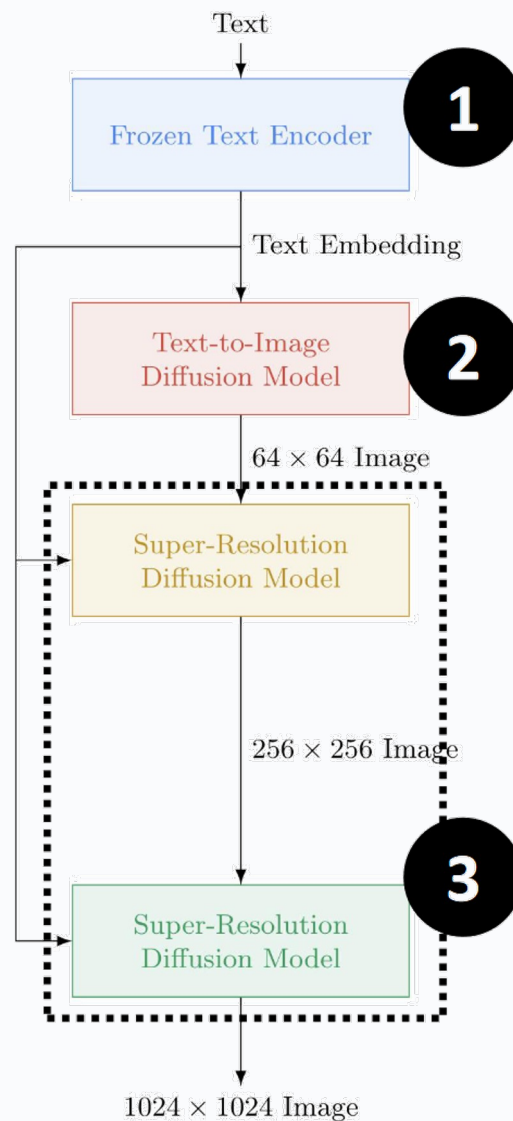



French Bread Cat Loaf Metal Print

Imagen

<https://arxiv.org/abs/2205.11487>

Google, 2022




“A Golden Retriever dog wearing a blue checkered beret and red dotted turtleneck.”



Saharia C, Chan W, Saxena S, et al. Photorealistic text-to-image diffusion models with deep language understanding[J]. Advances in Neural Information Processing Systems, 2022, 35: 36479-36494.

DALL-E2, OpenAI, 2022



图像和文本编码在同一个特征空间的好处




a photo of a cat → an anime drawing of a super saiyan cat, artstation



a photo of a victorian house → a photo of a modern house



a photo of an adult lion → a photo of lion cub



a photo of a landscape in winter → a photo of a landscape in fall



Diffusion Transformer

<https://arxiv.org/abs/2212.09748>



去噪：

输入：带噪声的Latent

输出：噪声的预测



DiT, 2022(2023)
<https://arxiv.org/abs/2212.09748>



Flow Matching

Lipman Y, Chen R T Q, Ben-Hamu H, et al. Flow matching for generative modeling[J]. arXiv preprint arXiv:2210.02747, 2022.
Liu X, Gong C, Liu Q. Flow straight and fast: Learning to generate and transfer data with rectified flow[J]. arXiv preprint arXiv:2209.03003, 2022.
Albergo M S, Vanden-Eijnden E. Building normalizing flows with stochastic interpolants[J]. arXiv preprint arXiv:2209.15571, 2022.

Flow-based Model as SCALE



Generated by MovieGen, Meta, 2024

A girl is running across a beach and holding a kite. She's wearing jean shorts and a yellow t-shirt. The sun is shining down.



Generated by Flux, Black Forest Labs, 2024

Flow-based Model as SCALE




Generated by Flux, Black Forest Labs, 2024



Generated by Flux, Black Forest Labs, 2024


生成模型的本质

VAE, GAN, Diffusion, Flow, etc.




扩散模型的两个局限

- 如果初始分布不是Gaussian?
- 如何更快地生成样本?



FM的工作原理

你需要先决定走多久停下来看一看



FM的训练：走直线

训练

